

Digitization of the Australian Parliamentary Debates, 1998-2022

Lindsay Katz*

Rohan Alexander†

January 20, 2023

Public knowledge of what is said in parliament is a key tenet of democracy, and a critical resource for political science research. In Australia, following the British tradition, the written record of what is said in parliament is known as Hansard. While the Australian Hansard has always been publicly available, it has been difficult to use it for the purpose of large-scale macro and micro-level text analysis because it has not been available as a dataset of sufficient quality. Following the lead of the Linked Parliamentary Data project which achieved this for Canada, we provide a new, comprehensive, high-quality database that captures all proceedings of the Australian parliamentary debates from 1998 to the present using Hansard. The database is publicly available and can be readily linked to other datasets such as election results. The creation and accessibility of this database will enable the exploration of questions that are not currently possible to explore, serving as a valuable resource for both researchers and policymakers.

1 Background & Summary

The official written record of parliamentary debates, formally known as Hansard, plays a fundamental role in capturing the history of political proceedings and facilitating the exploration of valuable research questions. Originating in the British parliament, the production of Hansard became tradition in a number of Commonwealth countries such as Canada, the United Kingdom, and Australia (Vice and Farrell 2017). Given the content and magnitude of these records, they have significance, particularly in the context of political science research. In the case of Canada, a team of researchers at the University of Toronto have digitized Hansard from 1901 to 2019, an endeavour called the Linked Parliamentary Data (LiPaD) project (Beelen et al. 2017). Having a digitized version of Hansard enables researchers to perform advanced analyses

*University of Toronto, lindsay.katz@mail.utoronto.ca

†University of Toronto, rohan.alexander@utoronto.ca

on these records using text analysis tools and statistical modelling. Following the lead of the LiPaD project, in this paper we introduce a novel database for the Australian Hansard. This is composed of individual datasets for each sitting day in the House of Representatives from 1998 to 2022, containing details on everything said in parliament in a form which can be readily used by researchers. With the development of tools for large-scale text analysis, this database will serve as a valuable resource for efficiently studying and exploring political behaviour in Australia over time.

The House of Representatives performs a number of crucial governmental functions, such as creating new laws and overseeing government expenditure (House of Representatives 2018, ch. 1). The House operates under a parallel chamber setup, meaning there are two debate venues where proceedings take place: the Chamber, and the Federation Chamber. Sittings of the House follow a predefined order of business, regulated by procedural rules called standing orders (House of Representatives 2018, ch. 8). A typical sitting day in the Chamber has a number of scheduled proceedings including debates on government business, 90 second member statements, and Question Time (House of Representatives 2018, ch. 8). The Federation Chamber was created in 1994 as a subordinate debate venue of the Chamber. This allows for better time management of House business as its proceedings occur simultaneously with those of the Chamber (House of Representatives 2018, ch. 21). Sittings in the Federation Chamber are quite different to those of the Chamber in terms of their order of business and scope of discussion. Business matters discussed in the Federation Chamber are limited largely to intermediate stages of bill development, and the business of private Members (House of Representatives 2018, ch. 21). It is the recording and compilation of all of these proceedings on which Hansard is based, and it is essentially, but not entirely, verbatim.

A week or so after each sitting day, a transcript is available for download from the official Parliament of Australia website in both static PDF and extensible markup language (XML) form. The PDF is the official release, which is converted to typed text using Optical Character Recognition (OCR) technology (Sherratt 2016). This conversion allows for the production of the XML formatted transcript, though these are not always perfect conversions and are particularly flawed in the early decades of Hansard (Sherratt 2016). The PDF imposes formatting designed for humans to read with ease, whereas XML is designed for consistency and machine legibility. The nature of XML enables us to more easily use code to manipulate these records at scale, motivating our choice to develop our database solely using the XML formatted files. In cases where we were unsure on how to proceed with processing the XML, we defer first to the PDF, and then to the video recording of the proceeding if available.

At present, the Hansard format that is available on the Parliament of Australia website is not accessible for large scale analysis. To this point, various researchers have had to create their own databases in an attempt to produce usable, complete data based on content from the Australian Parliament website. For instance, Sherratt (2016) created an online, easy to read database of Hansard from 1901 to 1980 using the XML files. These data can be navigated by year, parliament, people, and bills (Sherratt 2016). In an effort to make the Australian Parliamentary Handbook more accessible, Leslie (2021) has released an R package

which includes data on all Members of the Australian House of Representatives from 1945 to 2019. Further, Alexander and Hodgetts (2021) created the **AustralianPoliticians** R package, which contains a number of datasets related to the political and biographical information of Australian federal politicians who were active between 1901 and 2021.

A number of papers exist which utilize components of Australian Hansard to explore various research topics. For example, Salisbury (2011) used online Hansard records to investigate the occurrences of unparliamentary comments by members of the House, where The Speaker tells that member to withdraw their remark. Rasiah (2010) worked with Question Time data from Hansard transcripts during February and March of 2003, to investigate resistance of politicians in answering questions about Iraq. In their work, Fraussen, Graham, and Halpin (2018) use Hansard to quantify political prominence by investigating strategic mentions of interest groups by elected officials. Finally, the work of Alexander and Alexander (2021) offers a dataset of the Australian Hansard, along with an analysis of the effect of elections and changes in Prime Ministers upon topics mentioned in parliament. Alexander and Alexander (2021) created this database with the static PDF versions of Hansard, using OCR to digitize these files into text which is suitable for analysis. This means there are considerable digitization errors especially in the first half of the dataset.

While there is evidently a growing body of literature on this topic, there is still no comprehensive database for Australian Hansard based on XML that spans from 1901 to the present day. Our work serves to bridge this gap.

2 Methods

Our database contains one comma-separated value (CSV) file for each sitting day of the House of Representatives from 02 March 1998 to 08 September 2022. We developed four scripts to produce all of these files, where each script is designed to parse Hansard documents from a specific portion of the 1998 to 2022 time frame.

This section will be structured as follows. First, we will provide an overview of our approach to understanding and parsing an individual Hansard XML document, which informed all four scripts used to create our database. This will be supplemented with an excerpt from a Hansard XML to provide a visual example of its structure. Next we will explain the specific differences between the four scripts, and outline what structural changes necessitated their separate development. We will then provide details on the methodological intricacies of three core components of Hansard proceedings: Question Time, interjections, and stage directions.

2.1 Overview

The approach to parsing contents of an XML document is heavily dependent on its tree structure. As such, to create this database, we started by looking at a single Hansard XML

transcript from 2019. Doing so enabled us to identify the various components of interest in the document, and how each one can be parsed according to its corresponding structural form. Parsing was performed in R using the XML and xml2 packages (Temple Lang 2022; Wickham, Hester, and Ooms 2021). Focusing on one transcript also allowed us to ensure that all key components of the transcript were being parsed and captured with as much detail as possible. The typical form of a Hansard XML transcript is summarized in the nested list below. Note that this provides a broad overview of the structure, and does not contain every possible nested element that may be found in a Hansard XML.

```

<hansard>
  1. <session.header>
  2. <chamber.xscript>
    a) <business.start>
    b) <debate>
      i. <debateinfo>
      ii. <debate.text>
      iii. <speech>
      iv. <subdebate.1>
        (1) <subdebateinfo>
        (2) <subdebate.text>
        (3) <speech>
        (4) <subdebate.2>
          (a) <subdebateinfo>
          (b) <subdebate.text>
          (c) <speech>
  3. <fedchamb.xscript>
  4. <answers.to.questions>
    a) <question>
    b) <answer>

```

The outer-most node, also known as the parent node, is denoted `<hansard>` and serves as a container for the entire document. This parent node may have up to four child-nodes, where the first child-node contains details on the specific sitting day. Next, `<chamber.xscript>` contains all proceedings of the Chamber, `<fedchamb.xscript>` contains all proceedings of the Federation Chamber, and `<answers.to.questions>` contains Question Time proceedings. Note that the Federation Chamber does not meet on every sitting day, so this child element is not present in every XML file. The use of separate child nodes allows for the distinction of proceedings of the Chamber and Federation Chamber. The structure of the `<chamber.xscript>` and `<fedchamb.xscript>` nodes are generally the same, where the proceeding begins with a business start which is followed by a series of debates. Debate nodes can contain a sub-debate 1 child-node which has a sub-debate 2 child-node nested within it. That said, sometimes sub-debate 2 is not nested within sub-debate 1. Each of these three elements (i.e. `<debate>`, `<subdebate.1>`, and `<subdebate.2>`) as well as their respective sub-elements contain impor-

tant information on the topic of discussion, who is speaking, and what is being said. The `<speech>` node within each one contains the bulk of the text associated with that debate or sub-debate. A typical `<speech>` node begins with a `<talk.start>` sub-node, providing information on the speaker and the time of their first statement. Unsurprisingly, speeches rarely go uninterrupted in parliamentary debate settings — they are often composed of a series of interjections and continuations. These statements are categorized under different sub-nodes depending on their nature, such as `<interjection>` or `<continuation>`. The final key component of Hansard is Question Time, in which questions and answers are classified as unique elements. More detail on the purpose and processing of Question Time will follow in Section 2.3.

Figure 1 provides an example of the beginning of an XML file for Hansard, which illustrates the structure outlined in the nested list above. As stated, the XML structure begins with a parent element "hansard" (highlighted in blue), followed by a child element "session.header" (highlighted in yellow) with sub-child elements such as the date and parliament number, which are all highlighted in pink. Next there is the child element containing everything that takes place in the Chamber, "chamber.xscript", which is also highlighted in yellow in Figure 1. As previously mentioned, the first sub-node of "chamber.xscript" is the business start. The structure of the business start can be seen between the nodes highlighted in green in Figure 1, where the content we parse from the business start are highlighted in orange.

does this feel a bit repetitive? do you think it's worth keeping?

```
<hansard noNamespaceSchemaLocation="../../../hansard.xsd" version="2.2">
  <session.header>
    <date>2020-02-25</date>
    <parliament.no>46</parliament.no>
    <session.no>1</session.no>
    <period.no>2</period.no>
    <chamber>House of Reps</chamber>
    <page.no>0</page.no>
    <proof>0</proof>
  </session.header>
  <chamber.xscript>
    <business.start>
      <body background="" style="" xmlns:w="http://schemas.openxmlformats.org/wordprocessingml/2006/main" xmlns:a="http://schemas.openxmlformats.org/drawingml/2006/main" xmlns:o="urn:schemas-microsoft-com:office:office" xmlns:v="urn:schemas-microsoft-com:vml" xmlns:W="http://schemas.microsoft.com/office/word/2003/auxHint" xmlns:aml="http://schemas.microsoft.com/aml/2001/core" xmlns:pic="http://schemas.openxmlformats.org/drawingml/2006/picture" xmlns:w10="urn:schemas-microsoft-com:office:word" xmlns:wp="http://schemas.openxmlformats.org/drawingml/2006/wordprocessingDrawing" xmlns:r="http://schemas.openxmlformats.org/officeDocument/2006/relationships">
        <p class="HPS-SODJobDate" style="direction:ltr;unicode-bidi:normal;">
          <span class="HPS-SODJobDate">
            <span style="font-weight:bold;" />
            <a href="Chamber" type="">Tuesday, 25 February 2020</a>
          </span>
        </p>
        <p class="HPS-Normal" style="direction:ltr;unicode-bidi:normal;">
          <span class="HPS-Normal">
            <span style="font-weight:bold;">The SPEAKER</span>
            <span style="font-weight:bold;">Hon.</span>
          </span>
          <span style="font-weight:bold;">Tony Smith</span>
          <span style="font-weight:bold;">took the chair at 12:00, made an acknowledgement of country and read prayers.</span>
        </p>
      </body>
    </business.start>
  </chamber.xscript>
</hansard>
```

Figure 1: Snapshot of beginning of XML file for Hansard on 25 February 2020

Evidently, the nature of XML formatting means that different pieces of information are categorized under a series of uniquely named and nested nodes. As a result, to parse each piece of information, one must specify the unique hierarchy of nodes in which it is structured. This

is known as an XPath expression, and tells the parser how to navigate the XML document to obtain the desired information. For example, the session header date in Figure 1 can be accessed using the XPath expression “`hansard/session.header/date`”. We began our first script by parsing every piece of information contained in the XML document, using these unique XPath expressions to do so.

The next step was to further develop our script to produce tidy data sets containing each parsed element, where each statement is separated onto its own row with details about the speaker, and rows are placed in chronological order. This first involved correcting the variable classes and adding a number of indicator variables to differentiate where statements came from, such as Chamber versus Federation Chamber or sub-debate 1 versus sub-debate 2. The next key task stemmed from the fact that the raw text data were not separated by each statement when parsed. In other words, any interjections, comments made by The Speaker or Deputy Speaker and continuations within an individual speech were all parsed together as a single string. As such, the name, name ID, electorate and party details were only provided for the person whose turn it was to speak. Splitting up these speeches in a way which would be generalizable across sitting days required much thought and effort. Section 2.4 will provide further details on the intricacies of this task.

Since we are looking at a wide time span of documents, there are expectedly a number of changes in the way they are formatted over time. These became apparent as we ran our script on XML files from earlier sitting days. Some changes are as subtle as a differently named child-node, while others are as extensive as a completely different nesting structure. Smaller changes were accounted for as we became aware of them, and embedded into the code in a way that would not cause issues for parsing more current Hansards with subtle differences in formatting. However, as mentioned, more significant changes in the XML structure of Hansard are what necessitated the development of separate scripts as we worked backwards in time. Further, while working backward in time, it became clear that not every sitting day contains every possible XML element. For example, some days did not have sub-debate 2 content, and some days did not have a Federation Chamber proceeding. To improve the generalizability of these scripts, if-else statements were embedded within the code wherever an error might arise due to a missing element. For example, the entire Federation Chamber block of code is wrapped in an if-else statement for each script, so that it only executes if what the code attempts to parse actually exists in the file.

Once the script ran without error for a few recent years of Hansard, we continued to work backwards until extensive changes in tree structure made our script incompatible with parsing earlier XML files. The earliest sitting day our first script can successfully parse is 14 August 2012. Before developing new scripts to parse earlier Hansard documents, we decided to prioritize cleaning and finalizing what we have been able to parse. As such we continued building our script, fixing any problems we noticed in the resulting datasets such as excess whitespace or spacing issues, and splitting up any additional sections of the parsed text onto separate rows where necessary. Specifically, we added a section of our script to separate out general stage directions. More information on this separation will be provided in Section 2.5.

After completing our first script, it was formatted as a function which takes a single file name argument and produces one CSV file containing data on all proceedings from the given sitting day.

2.2 Script Differences

As previously mentioned, we developed a total of four scripts to parse the 1998-2022 time frame of Hansard documents. Two main factors motivated us to create four scripts as opposed to just one, the first being structural variation in XML over time, and the second being improved computational efficiency with separate scripts. While all four scripts use the same general approach to parsing described in Section 2.1 and produce the same CSV structure, the first and second scripts use a different method of data processing than the third and fourth scripts.

The need for a second script stems from the fact that when established in 1994, the Federation Chamber was originally named the Main Committee. The Main Committee was renamed to the Federation Chamber in mid-2012 (House of Representatives 2018, ch. 21). As a result, the child-node under which Federation Chamber proceedings are nested is named `<maincomm.xscript>` in all XML files prior to 14 August 2012. Having developed our first script based on Hansard from recent years, all XPath expressions contain the `<fedchamb.xscript>` specification. To avoid causing issues in our first script which successfully parses about 10 years of Hansard, we created a second script where we replaced all occurrences of `<fedchamb.xscript>` with `<maincomm.xscript>`. After making this modification and accounting for other small changes such as time stamp formatting, this second script successfully parses all Hansard sitting days from 10 May 2011 to 28 June 2012 (inclusive).

While the modifications needed to develop the second script were quite straightforward, this was not the case for our next script. The typical tree structure of Hansard XMLs spanning from 1998 to March 2011 has an important difference from that of XMLs released after March 2011, necessitating many changes to be made in our methodology. In XMLs after March 2011, which our first two scripts successfully parse, the first two child-nodes of `<speech>` are typically `<talk.start>`, and `<talk.text>`. The first child-node contains data on the person whose turn it is to speak, and the second contains the entire contents of that speech — including all interjections, comments and continuations. After the `<talk.text>` element closes, there are typically a series of other child-nodes which provide a skeleton structure for how the speech proceedings went in chronological order. For example, if the speech began, was interrupted by one member, and then continued uninterrupted until the end, there would be one `<interjection>` node and one `<continuation>` node following the `<talk.text>` node. These would contain details on the person who made each statement such as their party and electorate.

In contrast, the speech contents in XMLs from 1998 up to and including 24 March 2011 are nested differently — there is no `<talk.text>` node. Rather than this single child-node that

contains all speech content, statements are categorized in individual child-nodes. This means that unlike our code for parsing more current Hansards, we cannot specify a single XPath expression such as “`debate/speech/talk.text`” or “`debate/subdebate.1/speech/talk.text`” to extract all speeches, in their entirety, at once. This difference in nesting structure made our second script unusable for parsing transcripts preceding 10 May 2011, and required us to change our data processing approach considerably.

Since there are numerous speeches in a typical Hansard, each having multiple child-nodes with the same names (i.e. `<continuation>`, `<interjection>`), it is difficult to parse everything while maintaining the correct ordering of statements. We found that the most straightforward way to preserve the ordering of statements and to parse all speech contents at once was to parse from the `<debate>` element directly. The reason we did not use its `<speech>` child-node is because every speech has a unique structure of node children, and this makes it difficult to write code for data cleaning which is generalizable across all speeches and sitting days.

The challenge with parsing through the `<debate>` element is that every piece of data stored in that element is parsed as a single string, including all `<talk.start>` data, and all nested sub-debate data. For example, the data shown in Figure 2 would be parsed as a single string preceding the speech content, like so:

09:31:0010261Costello, Peter, MPMr COSTELLOCT4HigginsLPTreasurer10

```
<speech>
  <talk.start>
    <talker>
      <time.stamp>09:31:00</time.stamp>
      <page.no>10261</page.no>
      <name role="metadata">Costello, Peter, MP</name>
      <name role="display">Mr COSTELLO</name>
      <name.id>CT4</name.id>
      <electorate>Higgins</electorate>
      <party>LP</party>
      <role>Treasurer</role>
      <in.gov>1</in.gov>
      <first.speech>0</first.speech>
    </talker>
    <para>-I move:</para>
  </talk.start>
```

Figure 2: Portion of XML file for Hansard on 12 December 2002

This was not isolated to just the beginning of speeches — details on individuals interjecting or commenting during speeches were also captured this way. To separate statements out correctly, we collected all of these patterns using the `<talk.start>` node, and used them to split statements wherever one of these patterns was found. After separating statements, we were later able remove these patterns from the body of text. We also used this method of extracting and later removing unwanted patterns for other pieces of data which did not belong to the debate proceedings, such as sub-debate titles.

Once we finalized this new method of processing the data, we proceeded with data cleaning using the same general approach as in the first two scripts to produce the same structure

of CSV output. We then worked backwards in time and modified the code as needed for generalizability. Throughout this process we found a number of transcription errors present in the XMLs from earlier years. We fixed these manually, deferring to the official release to ensure the correct information was filled in. Since there were a number of transcription errors specific to the 2000s, we chose to create a fourth script for parsing 1998 and 1999. This allowed us to remove all the code which was needed to resolve specific transcription errors of the 2000s, to avoid an overly long script and in turn improving computational efficiency. As such, our fourth script is essentially the same as the third, the only difference being that it has code specific to fixing transcription errors from 1998 and 1999.

2.3 Question Time

A key characteristic of the Australian parliament system is the ability for the executive government to be held accountable for their decisions. One core mechanism by which this is achieved is called Question Time. This is a period on each sitting day in the Chamber where members of the House can ask ministers two types of questions: questions in writing which are written in advance, or questions without notice which are asked verbally in the Chamber and are responded to in real time (House of Representatives 2021). Parsing the components of Question Time required a slightly different approach than that of the debate speech, because of its unique structure in the XML document. Questions without notice are included directly in the `"chamber.xscript"` child node, with sub-child nodes called `"question"` and `"answer"` to differentiate the two. Questions in writing, however, are embedded in their own child node called `"answers.to.questions"` at the end of the XML file.

To parse questions without notice and questions in writing in the first two scripts, we extracted all question and answer elements from the `"chamber.xscript"` and `"answers.to.questions"` child nodes. For the third and fourth scripts, our method of parsing the `"chamber.xscript"` speeches meant that all questions without notice content was already parsed in order. Similar to the rest of the speech content, we used those patterns of data preceding the text to separate questions and answers. Finally, since questions in writing exist in their own child node we were able to use the same parsing method for the third and fourth scripts as we did for the first and second scripts.

We then added binary flags to differentiate between questions and answers, determined by which sub-child node the statement was extracted from, that is, whether it came from a `<question>` node, or an `<answer>` node. Sometimes, questions were incorrectly transcribed under an answer node and vice-versa, in which cases we manually corrected the question and answer flags using the `str_detect` function from the `stringr` package. For instance, our code for parsing questions searches for any statements which include “has provided the following answer to the honourable member’s question”, in which case we re-code that statement as an answer. Once all questions and answers were correctly classified as such, we could then merge all the content back together such that each question was followed by its associated answer.

This was straightforward due to the fact that by nature of XML parsing, all questions and answers were parsed in the order in which they appeared in the transcript.

The next step was to merge Question Time contents with all the debate speech. Doing so with the questions without notice content was straightforward. For the first and second scripts, we used the time stamp and page number associated with each question to maintain the correct order of statements. For the third and fourth scripts, our method of parsing meant that everything was already parsed together in order, so we did not have to perform any additional merging. While questions in writing do not contain time stamps, merging this content was also straightforward due to the fact that it is always at the very end of Hansard. This means that we could simply bind question in writing rows to the bottom of the main dataframe. This approach was used for all four scripts.

2.4 Interjections

As mentioned, the text was structured and parsed in such a way that various interjections and comments which happened during a speech were not separated onto individual rows. This was the case across the entire time frame of documents. We will first discuss the methodology employed to split interjections in the first and second scripts, as it informed our approach for the third and fourth scripts.

Below is an example of part of a speech we would need to split, extracted from Hansard on 30 November 2021, where Mr. van Manen is interrupted by The Speaker who states that the time for members' statements has concluded.

“Mr VAN MANEN (Forde—Chief Government Whip) (13:59): It’s a great pleasure to share with the House that Windaroo Valley State High School has qualified for the finals of the Australian Space Design Competition, to begin in January next year. The competition is regarded as the premier STEM competition for high school students and is recognised by universities around the country. The students are required to respond to industry-level engineering and requests for tender for design and—The SPEAKER: Order! In accordance with standing order 43, the time for members’ statements has concluded.”

We want each statement on its own row with the correct name, name ID, electorate and party information on the individual speaking. We approached this task in a number of steps.

Once all parsed text from the XML was merged into one dataframe called `main`, our first step was to add a `"speech_no"` variable. This was done to keep track of which speech each interjection, comment, or continuation belonged to as we separated these components onto their own rows.

The next step was to extract all the names and titles preceding these interjections, comments and continuations. This would enable us to then separate the speeches in the correct places

using these names and titles in combination with regular expressions, which are patterns of characters that can be used to search bodies of text. We completed this extraction process with a few intermediate steps, due to the large number of name styles and interjection types that had to be accounted for, each requiring their own unique regular expression format.

As mentioned in Section 2.2, more recent years of Hansard XMLs contain a series of child-nodes which exist to capture the structure of interruptions in that speech. Figure 3 provides an example of this, where the speech was interrupted by a comment from The Deputy Speaker, and then the member continued their speech. Looking at the element names highlighted in blue, it is clear that these child-nodes do not contain the actual text for the interjection or continuation — this text is embedded within the speech above it. However, as shown by the content highlighted in pink in Figure 3, we were able to extract useful details on the individual interjecting which we could use later on. Making use of this structure, we extracted names and information of all individuals that were categorized within the XML as interjections. We stored this as a dataframe called "interject".¹

```

<interjection>
  <talk.start>
    <talker>
      <page.no>264</page.no>
      <time.stamp />
      <name role="metadata">O'Brien, Llew (The DEPUTY SPEAKER)</name>
      <name.id>10000</name.id>
      <electorate>Wide Bay</electorate>
      <party>LNP</party>
      <in.gov />
      <first.speech />
    </talker>
  </talk.start>
  <talk.text>
</interjection>
<continue>
  <talk.start>
    <talker>
      <page.no>264</page.no>
      <time.stamp />
      <name role="metadata">Mitchell, Rob, MP</name>
      <name.id>M3E</name.id>
      <electorate>McEwen</electorate>
      <party>ALP</party>
      <in.gov />
      <first.speech />
    </talker>
  </talk.start>
  <talk.text>
</continue>

```

Figure 3: Snapshot of XML structure with interjection and continuation from 03 February 2021 Hansard

We then created lists using both the `interject` and `main` dataframes to capture all the names of individuals who spoke that day. We added the names of all Members in a number of unique formats, due to the frequent variation in how names are transcribed in Hansard. When a Member interjects or continues a speech, the usual form of their name is a title followed by

¹We decided not to include this data in our final database, as it is all embedded in our resulting datasets which have a flag for interjections.

their first name or first initial and/or last name. There is also variation in the capitalization of these names.² Another source of variation is in individuals with more than one first name, as sometimes only their first first name is written, while other times their entire first name is written. Additionally, some surnames have punctuation, and some surnames have specific capitalization such as “McCormack”, where even in full capitalization, the first “c” remains lower case. This variation demands careful consideration when writing regular expression patterns to search for in the text. In these lists we also accounted for any general interjection statements transcribed that were not attributed to an individual, such as “An opposition member interjecting-”.

Having these lists enabled us to extract the names of Members and their associated titles as they exist in the text, by searching for exact matches with regular expression patterns. We then used these extracted names to split all the speeches up, using a number of regular expressions with lookarounds. A lookahead can be added to a regular expression pattern to enhance the specificity of matches. These were used to ensure that the text was not being split in the wrong places, such as places where Members were simply being named in the statement of another Member.

Once all interjections, comments and continuations were successfully split onto their own rows using the lists we created, we did one final check for any additional names that were not captured in these lists. To do so, we searched for any remaining name matches in speech bodies with general regular expressions and lookarounds, and separated text using those matches when found.

We then added an order variable to the dataset based on row number, to keep track of the order in which everything was said. The next step was to fill the name, name ID, electorate and party variables with the correct data for each row. We also wanted to add the gender and unique identifier for each individual as found in the **AustralianPoliticians** package. To do so, we created a lookup table, which contained the unique incomplete form in which the name was transcribed, and the corresponding full name, name ID, electorate, party, gender and unique ID for that individual. Figure 4 provides an example of this. We used the main dataset from the **AustralianPoliticians** package in the creation of each lookup table (Alexander and Hodgetts 2021).

name	gender	uniqueID	name.id_use	electorate_use	party_use	name_use
Mr RICK WILSON	male	Wilson1966	198084	O'Connor	LP	Wilson, Rick, MP
Mr Rob Mitchell	male	MitchellRobert1967	M3E	McEwen	ALP	Mitchell, Rob, MP
Ms JULIE BISHOP	female	Bishop1956	83P	Curtin	LP	Bishop, Julie, MP
The DEPUTY SPEAKER (Ms Vamvakinou)	female	Vamvakinou1959	00AMT	Calwell	ALP	Vamvakinou, Maria (The DEPUTY SPEAKER)
Ms Macklin	female	Macklin1953	PG6	Jagajaga	ALP	Macklin, Jenny, MP
Mr Keogh	male	Keogh1981	249147	Burt	ALP	Keogh, Matt, MP
Mr BUCHHOLZ	male	Buchholz1968	230531	Wright	LP	Buchholz, Scott, MP
Mr TURNBULL	male	Turnbull1954	885	Wentworth	LP	Turnbull, Malcolm, MP
Mr Turnbull	male	Turnbull1954	885	Wentworth	LP	Turnbull, Malcolm, MP

Figure 4: First 10 rows of lookup table from 19 October 2017 Hansard processing

²Sometimes when someones first name is included, only their last name is capitalized, while sometimes their full name is capitalized, or other times neither are capitalized.

Next, we merged our main dataframe with the lookup table to replace any incomplete names with their full names, and to fill in any gaps with available name ID, electorate, party, gender and unique ID information. Finally, we were able to add a flag for interjections. Grouping our data by the speech number, we defined an interjection as a statement made by anyone who is not The Speaker, The Deputy Speaker, or the Member whose turn it was to speak. Figure 5 provides an example of a Federation Chamber proceeding with interjections. Evidently, statements made by the original speaker Stuart Robert or by The Deputy Speaker Maria Vamvakinou are not flagged as interjections.

name	order	speech_no	page_no	time_stamp	name_id	electorate	party	body	fedchamb_flag	sub1_flag	sub2_flag	question	answer	q_in_writing	gender	uniqueID	interject	div_flag
Robert, Stuart, MP	371	181	3810	13:31:00	HWT	Fadden	LP	It is always a pleasure to follow the member for	1	0	0	0	0	0	male	Robert1970	0	0
An opposition member	372	181	3810	NA	NA	NA	NA	An opposition member interjecting-	1	1	0	0	0	0	NA	NA	1	0
Robert, Stuart, MP	373	181	3810	NA	HWT	Fadden	LP	It is 30 minutes! If the member were that pass!	1	0	0	0	0	0	male	Robert1970	0	0
Champion, Nick, MP	374	181	3810	NA	HWS	Wakefield	ALP	I have been a lot less successful.	1	1	0	0	0	0	male	Champion1972	1	0
Vamvakinou, Maria (The DEPUTY SPEAKER)	375	181	3810	NA	QOAMT	Calwell	ALP	Order! The level of interjection is too high.	1	1	0	0	0	0	female	Vamvakinou1959	0	0
Robert, Stuart, MP	376	181	3810	NA	HWT	Fadden	LP	Has the member for Wakefield ever seen som	1	0	0	0	0	0	male	Robert1970	0	0
Vamvakinou, Maria (The DEPUTY SPEAKER)	377	181	3810	NA	QOAMT	Calwell	ALP	The time allotted for this debate has expired. T	1	0	0	0	0	0	female	Vamvakinou1959	0	0
stage direction	378	181	3810	13:37:00	NA	NA	NA	Sitting suspended from 13:37 to 16:00	1	1	0	0	0	0	NA	NA	0	0

Figure 5: Example of speech with interjections from 21 November 2016 Hansard

Having developed a successful methodology for splitting interjections, we used this to inform our general approach in the third and fourth scripts. However, the difference in data cleaning used in these scripts necessitated some departure from the original methodology. As discussed in Section 2.2, we used string patterns extracted from `<talk.start>` nodes to separate speeches. As evident in Figure 3, `<talk.start>` nodes are nested within `<interjection>` nodes, meaning that the patterns of data from interjection statements were separated out in the process. In other words, our approach to data cleaning in the third and fourth scripts separated out interjections in the process. This meant that we did not need to create lists of names and titles to search for in the text as we did before. However, we used the same list of general interjection statements to separate on that was used in the first two scripts. We then did an additional check for statements that may have not been separated due to how they were embedded in the XML, and separated those out where needed.³

We then proceeded to clean up speeches and fill in correct speaker details. While we used the same lookup table approach as before, we did so in combination with another means of filling in speaker details. The patterns parsed from `<talk.start>` nodes contain important data on the speaker of each statement. As such, we could extract those data associated with each pattern by parsing one element inward, using the XPath expression `<talk.start/talker>`. We created a pattern lookup table with these data, and merged it with the main Hansard dataframe by the first pattern detected in each statement. Figure 6 provides an example of that lookup table. This approach enabled us to fill in missing data on each speaker using data extracted directly from the XML. Finally, we then used the `AustralianPoliticians` dataset to fill in other missing data, and flagged for interjections in the same manner as before.

³While most statements were categorized in their own child-node and hence captured through pattern-based separation, some were not individually categorized, and had to be split manually in this step.

page.no	time.stamp	name	name.short	name.id	electorate	party	role	in.gov	first.speech	first.pattern
10261	09:31:00	Costello, Peter, MP	Mr COSTELLO	CT4	Higgins	LP	Treasurer	1	0	09:31:0010261Costello, Peter, MPMr COSTELLOCT4Hi...
10262	09:36:00	Ruddock, Philip, MP	Mr RUDDOCK	OJ4	Berowra	LP	Minister for Immigration and Multicultural and Indige...	1	0	09:36:0010262Ruddock, Philip, MPMr RUDDOCKOJ4B...
10263	09:41:00	Williams, Daryl, MP	Mr WILLIAMS	7V5	Tangney	LP	Attorney-General	1	0	09:41:0010263Williams, Daryl, MPMr WILLIAMSV5Ta...
10264	09:47:00	Truss, Warren, MP	Mr TRUSS	GT4	Wide Bay	NP	Minister for Agriculture, Fisheries and Forestry	1	0	09:47:0010264Truss, Warren, MPMr TRUSSGT4Wide B...
10266	09:55:00	Truss, Warren, MP	Mr TRUSS	GT4	Wide Bay	NP	Minister for Agriculture, Fisheries and Forestry	1	0	09:55:0010266Truss, Warren, MPMr TRUSSGT4Wide B...
10267	10:00:00	Truss, Warren, MP	Mr TRUSS	GT4	Wide Bay	NP	Minister for Agriculture, Fisheries and Forestry	1	0	10:00:0010267Truss, Warren, MPMr TRUSSGT4Wide B...
10269	10:07:00	Truss, Warren, MP	Mr TRUSS	GT4	Wide Bay	NP	Minister for Agriculture, Fisheries and Forestry	1	0	10:07:0010269Truss, Warren, MPMr TRUSSGT4Wide B...
10269	10:09:00	Truss, Warren, MP	Mr TRUSS	GT4	Wide Bay	NP	Minister for Agriculture, Fisheries and Forestry	1	0	10:09:0010269Truss, Warren, MPMr TRUSSGT4Wide B...
10269	10:10:00	Slipper, Peter, MP	Mr SLIPPER	OV5	Fisher	LP	Parliamentary Secretary to the Minister for Finance an...	1	0	10:10:0010269Slipper, Peter, MPMr SLIPPEROV5Fisher...
10270	10:15:00	Slipper, Peter, MP	Mr SLIPPER	OV5	Fisher	LP	Parliamentary Secretary to the Minister for Finance an...	1	0	10:15:0010270Slipper, Peter, MPMr SLIPPEROV5Fisher...
10271	10:16:00	Worth, Trish, MP	Ms WORTH	8V5	Adelaide	LP	Parliamentary Secretary to the Minister for Health and...	1	0	10:16:0010271Worth, Trish, MPMs WORTH8V5Adelai...

Figure 6: 10 rows of pattern lookup table from 12 December 2012 Hansard processing

2.5 Stage Directions

When building our first scripts, one of the final components needed was to separate general stage directions out from statements made by members. Stage directions are general statements included in the transcript to document happenings in parliament. Examples of stage directions are “Bill read a second time”, “Question agreed to” or “Debate adjourned”. It was unclear to us from the XML and PDF who exactly these statements were attributed to. For further clarification, we watched portions of the video recording for some sitting days, and noticed that where these statements are documented in Hansard, they are not explicitly stated in parliament. For example, when The Deputy Speaker says “The question is that the bill be now read a second time”, members of the House take a vote, and if the majority is in favour, they proceed reading the bill the second time. This vote and second reading is not explicitly transcribed, rather what is written is: “Question agreed to. Bill read a second time”. For this reason, we filled the name variable for these statements with “stage direction”. Note that stage directions were not flagged as interjections. These stage directions are not defined differently from the regular debate speech in the XML, meaning we had to manually create a list of stage directions to separate out of the speeches. We have built this list of stage directions as we worked backwards in parsing Hansard, and took the same approach across all four scripts.

3 Data Records

Our database contains 1517 CSV files, covering all sitting days of the House of Representatives from 02 March 1998 to 08 September 2022. All data records are available on the general-purpose repository Zenodo, at <https://doi.org/10.5281/zenodo.7336076> (Katz and Alexander 2022). For each CSV file, each row contains an individual statement, with details on the individual speaking. Of course, for general statements transcribed as made by “Honourable members” for example, these variables cannot be specified. Table 1 provides an overview of each variable found in the database.

Table 1: Summary and description of variables in our database

Variable	Description
name	Name of speaker
order	Row number
speech_no	Speech number
page.no	Page number statement can be found on in official Hansard
time.stamp	Time of statement
name.id	Unique member identification code, based on the Parliamentary Handbook
electorate	Speaking member’s electorate
party	Speaking member’s party
in.gov	Flag for in government (1 if in government, 0 otherwise)
first.speech	Flag for first speech (1 if first speech, 0 otherwise)
body	Statement text
fedchamb_flag	Flag for Federation Chamber (1 if Federation Chamber, 0 if Chamber)
sub1_flag	Flag for sub-debate 1 contents (1 if sub-debate 1, 0 otherwise)
sub2_flag	Flag for sub-debate 2 contents (1 if sub-debate 2, 0 otherwise)
question	Flag for question (1 if question, 0 otherwise)
answer	Flag for answer (1 if answer, 0 otherwise)
q_in_writing	Flag for question in writing (1 if question in writing, 0 otherwise)
div_flag	Flag for division (1 if division, 0 otherwise)
gender	Gender of speaker
uniqueID	Unique identifier of speaker
interject	Flag for interjection (1 if statement is an interjection, 0 otherwise)

The `name`, `page.no`, `time.stamp`, `name.id`, `electorate`, `party`, `in.gov`, `first.speech`, and `body` variables all came directly from the XML contents. In addition to these variables, we added a number of flags to enable easy filtering of statements. For example, adding the `fedchamb_flag` provides a clear distinction between the proceedings of the Chamber with those of the Federation Chamber. As well, the `sub1_flag` and `sub2_flag` variables allow us to keep track of where various statements are being parsed from in the XML document. The `question`, `answer`, and `q_in_writing` flags were added to identify statements belonging to Question Time, and the nature of these statements. We also flagged for interjections (`interject`), and the `div_flag` variable was added to flag when a division is called for. The `gender` and `uniqueID` variables were added based on the main dataset from the `AustralianPoliticians` package. Details on the usage of `uniqueID` will be provided in Section 5. Further, the `speech_no` variable allows us to keep track of the speech number that each statement and interjection belongs to. Having the speech number variable offers an easy way to group statements by speech or isolate specific speeches of interest. Finally, the `order` variable was added to maintain the order of proceedings.

this paragraph is a bit repetitive and states everything found in the table, but i also feel like it's valuable to specify which ones we added and elaborate on their meaning. thoughts?

3.1 Descriptive Statistics

4 Technical Validation

We developed a script to perform seven automated tests on each file in our database, in an effort to enhance its quality and consistency. Our first test validates that the date specified in each file name matches the date specified in its corresponding XML session header.⁴ Every file passed this test, and we detected one discrepancy in an XML file from 03 June 2009, where its session header contained the wrong date. We validated that our file name and date was correct by checking the official PDF release from that sitting day.

The second test checks that in each file, there is an equal number of flagged questions and answers. This test was motivated by the discrepancies in correct categorization of questions and answers we found when parsing Question Time, and allows us to check for the presence of any of these discrepancies which we have missed.

When a member runs out of allotted time for their speech, Hansard editors transcribe “(Time expired)” after their final word. As a means of checking that we have separated speeches out correctly, our third test checks that when the phrase “(Time expired)” exists in a body of text, it exists at the very end. When this is not the case, we know that we have missed the separation of the next statement onto its own row, and could fix this accordingly.

The remaining tests focus on the members present on each sitting day. Our fourth test checks that there is one unique party and electorate attributed to each individual on each sitting day. As we parsed Hansard further back in time, we found a number of cases where an individual was associated with the wrong electorate or party due to transcription errors. When we found these data errors we corrected them based on the official release. This test provides us with an automated way to catch these errors and correct them at scale.

Next, we test that the unique name identification code attributed to each individual is found in the Australian Parliamentary Handbook. We do so using the `ausPH` package. This test serves as another means to correct for transcription errors, this time in the case of name IDs. We found and corrected for a number of common name ID transcription errors detected by this test, such as a capital letter “O” in place of a zero.

Our sixth test checks that on any given sitting day, the individuals identified are alive. To do so, we utilized the main dataset from the `AustralianPoliticians` package which contains

⁴As seen in Figure 1, the first child node of the `<session.header>` element is the date.

the birth and where applicable death dates for every politician. This test confirmed that all members who are detected to speak on each sitting day are not deceased.

Finally, our seventh test validates that all individuals speaking are active members of parliament on that particular day. We use the `mps` dataset from the `AustralianPoliticians` package which has the dates of each Member's duration in parliament. Using these dates, we check that each person speaking on each sitting day is in fact a Member of Parliament on that day.

5 Usage Notes

To enhance the usability of our database, we added a `uniqueID` variable to each CSV file. This serves as a unique identifier for each speaking Member, and comes from the `uniqueID` variable present within data from both the `AustralianPoliticians` and `AustralianElections` R packages (Alexander and Hodgetts 2021; Alexander 2019). By including this variable, one can readily integrate our data records with those available in these two packages.

Further, the `name.id` variable found in each CSV is another unique identifier for each Member of Parliament. This variable was parsed directly from the Hansard XML files, and can be found in the Australian Parliamentary Handbook. As such, our data records can be integrated with those from the `ausPH` package which provides datasets for contents of the Australian Parliamentary Handbook (Leslie 2021). This will allow for convenient extraction of further details on each Member of Parliament in a tidy, ready to analyze format.

6 Code Availability

The code written to build this database is available on the GitHub repository associated with this paper: <https://github.com/lindsaykatz/hansard-proj>. All scripts were created using R software (R Core Team 2022). The core packages used to develop these scripts are: the `XML` package by Temple Lang (2022), the `xml2` package by Wickham, Hester, and Ooms (2021), the `tidyverse` R packages by Wickham et al. (2019), the `AustralianPoliticians` package by Alexander and Hodgetts (2021), and the `ausPH` package by Leslie (2021). `XML` and `xml2` were used for parsing the XML documents, `AustralianPoliticians` and `ausPH` were used for cleaning up and filling in member details in the datasets, and `tidyverse` packages were used in all steps, for tidy wrangling of data.

7 Acknowledgements

8 Author Contributions

9 Competing Interests

The authors declare no competing interests.

References

- Alexander, Rohan. 2019. *AustralianElections: Datasets on Australian Elections*. <https://github.com/RohanAlexander/AustralianElections>.
- Alexander, Rohan, and Monica Alexander. 2021. “The Increased Effect of Elections and Changing Prime Ministers on Topics Discussed in the Australian Federal Parliament Between 1901 and 2018.” *arXiv Preprint arXiv:2111.09299*.
- Alexander, Rohan, and Paul A. Hodgetts. 2021. *AustralianPoliticians: Provides Datasets about Australian Politicians*. <https://CRAN.R-project.org/package=AustralianPoliticians>.
- Beelen, Kaspar, Timothy Alberdingk Thijm, Christopher Cochrane, Kees Halvemaan, Graeme Hirst, Michael Kimmins, Sander Lijbrink, et al. 2017. “Digitization of the Canadian Parliamentary Debates.” *Canadian Journal of Political Science/Revue Canadienne de Science Politique* 50 (3): 849–64.
- Fraussen, Bert, Timothy Graham, and Darren R Halpin. 2018. “Assessing the Prominence of Interest Groups in Parliament: A Supervised Machine Learning Approach.” *The Journal of Legislative Studies* 24 (4): 450–74.
- House of Representatives. 2018. Edited by D. R. Elder and P. E. Fowler. *House of Representatives Practice*. 7th ed. Australian Government - Department of the House of Representatives.
- . 2021. *A window on the house: Practices and procedures relating to question time*. Parliament of Australia.
- Katz, Lindsay, and Rohan Alexander. 2022. “A new, comprehensive database of all proceedings of the Australian Parliamentary Debates.” Zenodo. <https://doi.org/10.5281/zenodo.7336076>.
- Leslie, Patrick. 2021. *Members of the Australian House of Representatives (1945-2019)*. <https://github.com/palesl/AustralianHouseOfRepresentatives>.
- R Core Team. 2022. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Rasiah, Parameswary. 2010. “A Framework for the Systematic Analysis of Evasion in Parliamentary Discourse.” *Journal of Pragmatics* 42 (3): 664–80.

- Salisbury, Christopher. 2011. “’Mr Speaker, i Withdraw...’: Standards of (Mis) Behaviour in the Queensland, Western Australian and Commonwealth Parliaments Compared via Online Hansard.” *Australasian Parliamentary Review* 26 (1): 166–77.
- Sherratt, Tim. 2016. “Documentation: Historic Hansard.” *Historic Hansard*. <http://timsherratt.org/digital-heritage-handbook/docs/historic-hansard/#>.
- Temple Lang, Duncan. 2022. *XML: Tools for Parsing and Generating XML Within r and s-Plus*. <https://CRAN.R-project.org/package=XML>.
- Vice, John, and Stephen Farrell. 2017. *The History of Hansard*. House of Lords Library; House of Lords Hansard.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D’Agostino McGowan, Romain François, Garrett Golemund, et al. 2019. “Welcome to the tidyverse.” *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Wickham, Hadley, Jim Hester, and Jeroen Ooms. 2021. *Xml2: Parse XML*. <https://CRAN.R-project.org/package=xml2>.